Where subset-sum inputs fall in runtime

Saved measurements from this repository | logarithmic runtime axes

576 unscaled inputs from 144 lists; n = 24, 32, 38. No solver or model was rerun.

A Runtime histogram: frozen exact-reference portfolio

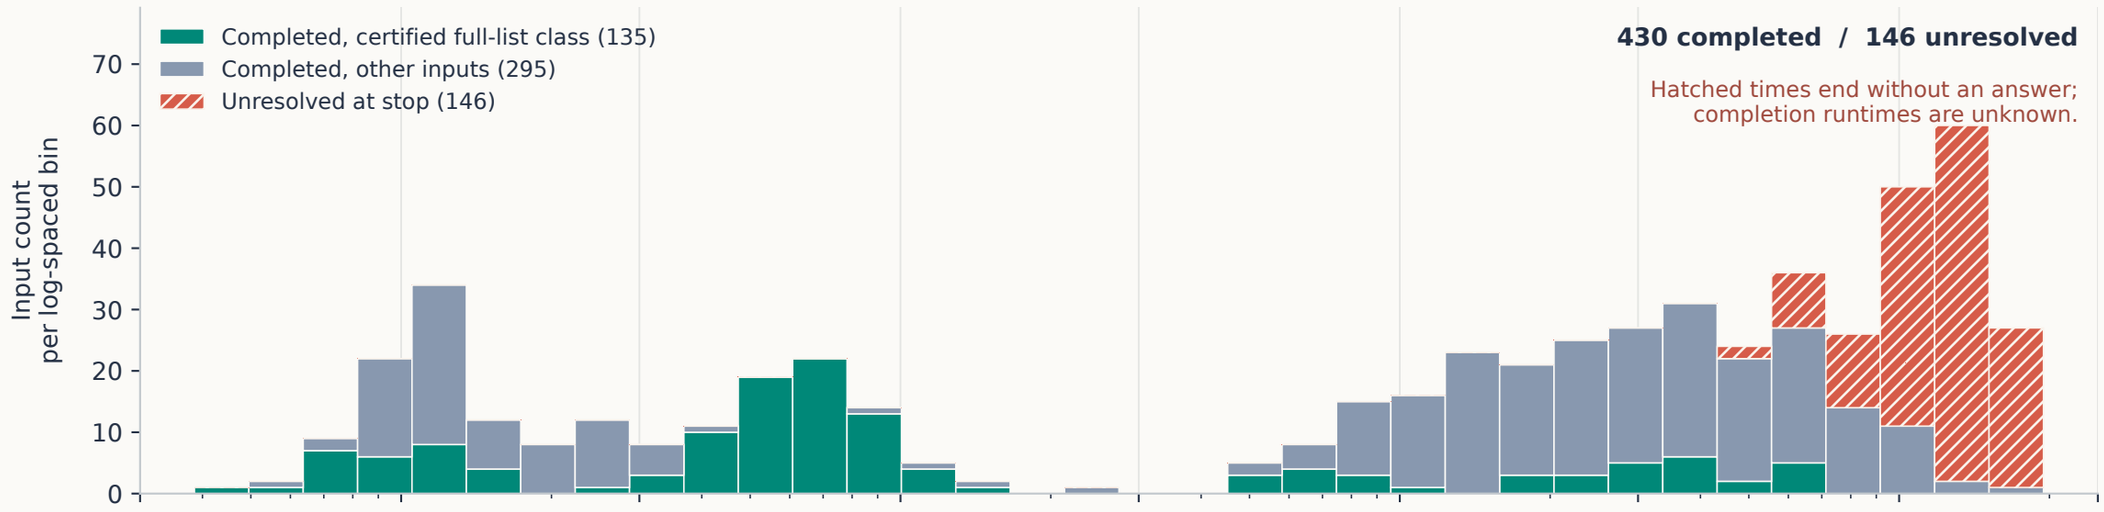

B Class locations and saved ML flags (groups overlap)

C Constructed hard NO cases: separate historical benchmark

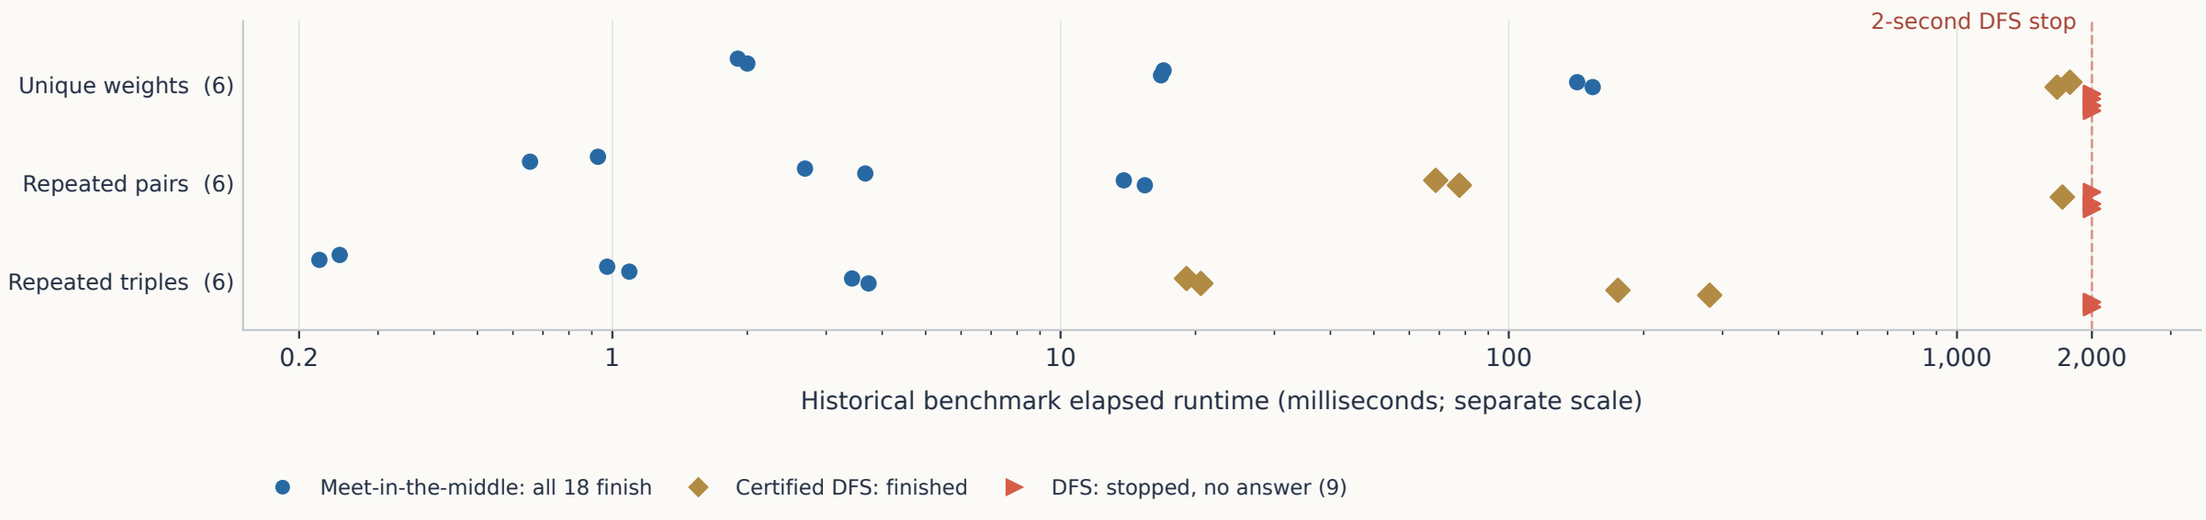

A–B: elapsed time of work-capped runs (8L allowance + state/bitset caps); includes exact preprocessing, excludes ML acquisition.
ML flag: original combined logistic model, p(completion) < 0.5; 118 unresolved + 23 completed flagged, 28 unresolved missed.
C: 41-bit odd weights, n = 24/30/36, two seeds; targets pass root bounds/residue checks. Different solvers/study; not pooled into A.
These are sampled synthetic inputs. Unresolved runs and ML flags do not establish intrinsic hardness or new complexity classes.